

Kaggle competition: Digit Recognizer

Learn computer vision fundamentals
with the famous MNIST data

Index

- Abstract
- Introduction
- Data Description
- Data Preparation
- Method – CNN
- Result
- Reference

Abstract


- This competition on Kaggle using MNIST dataset are for the beginners to better understand the basic of machine learning and application of neural networks by focusing on the classic topic "Handwritten Digit Recognition".
- MNIST ("Modified National Institute of Standards and Technology") is a database which has about 60,000 of digital handwriting image data, released in 1999 by Yann LeCun, Corinna Cortes and Christopher J.C. Burges.
- It was often used as a training data in classification models or other image processing algorithms.

Introduction

- In this competition, the goal is to correctly identify digits from a dataset of tens of thousands of handwritten images. We've curated a set of tutorial-style kernels which cover everything from regression to neural networks. We encourage you to experiment with different algorithms to learn first-hand what works well and how techniques compare.


Data Description

- The data files contain gray-scale images of hand-drawn digits, from zero through nine.
Each image is 28 pixels in height and 28 pixels in width, for a total of 784 pixels in total. Each pixel has a single pixel-value associated with it, indicating the lightness or darkness of that pixel, with higher numbers meaning darker. This pixel-value is an integer between 0 and 255, inclusive.
- Visually, if we omit the "pixel" prefix, the pixels make up the image like this:
 - 000 001 002 003 ... 026 027
 - 028 029 030 031 ... 054 055
 - 056 057 058 059 ... 082 083
 - | | | | ... | |
 - 728 729 730 731 ... 754 755
 - 756 757 758 759 ... 782 783



Data Preparation

- Data distribution
checking the counts of each digits.



```
1 4684
7 4401
3 4351
9 4188
2 4177
6 4137
0 4132
4 4072
8 4063
5 3795
Name: label, dtype: int64
```

Data Preparation

- Data cleaning
 - Check for null and missing values

```
# Check the data
```

```
X_train.isnull().any().describe()
```

```
Count    784  
Unique    1  
Top       False  
Freq      784  
dtype: object
```

```
test.isnull().any().describe()
```

```
Count    784  
Unique    1  
Top       False  
Freq      784  
dtype: object
```

Data Preparation

- Normalization

Normalize the dataset from
[0..255] to [0..1]

```
# Normalize the data  
X_train = X_train / 255.0  
test = test / 255.0
```


Data Preparation

- Reshape
reshaping all data to 28x28x1 3D matrices
due to Keras requires an extra dimension in
the end which correspond to channels.

```
# Reshape image in 3 dimensions (height = 28px, width = 28px , canal = 1)  
X_train = X_train.values.reshape(-1,28,28,1)  
test = test.values.reshape(-1,28,28,1)
```

Data Preparation

- One-hot encoding

Labels in MNIST dataset are 0 to 9, need to encode these numbers in to one hot vectors.

```
# Encode labels to one hot vectors (ex : 2 -> [0,0,1,0,0,0,0,0,0,0])  
Y_train = to_categorical(Y_train, num_classes = 10)
```

Data Preparation

- Split training and validation set
 split the training set into two parts:
 validation set (10%)
 training set (90%)

```
# Set the random seed
```


```
random_seed = 2
```

```
# Split the train and the validation set for the fitting
```

```
X_train, X_val, Y_train, Y_val = train_test_split(X_train, Y_train, test_size =  
0.1, random_state=random_seed)
```

Method – CNN

- CNN is used for image classification, object detection
- Using the Keras Sequential API here so we can just add one layer at a time, starting from the input.



Method – CNN

- Epoches and batch size

Each epoch takes about 10 mins to run on a windows laptop with a batch size of 86

```
epochs = 30  
batch_size = 86
```

Method – CNN

- Data augmentation
 - Without data augmentation, our accuracy is around 98.114%
 - With the data augmentation we reached the accuracy of 99.428%


Method – CNN

- Data augmentation

```
datagen = ImageDataGenerator(  
    featurewise_center=False, # set input mean to 0 over the dataset  
    samplewise_center=False, # set each sample mean to 0  
    featurewise_std_normalization=False, # divide inputs by std of the dataset  
    samplewise_std_normalization=False, # divide each input by its std  
    zca_whitening=False, # apply ZCA whitening  
    rotation_range=10, # randomly rotate images in the range (degrees, 0 to 180)  
    zoom_range = 0.1, # Randomly zoom image  
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total  
width)  
    height_shift_range=0.1, # randomly shift images vertically (fraction of total  
height)  
    horizontal_flip=False, # randomly flip images  
    vertical_flip=False) # randomly flip images  
  
datagen.fit(X_train)
```



Result

- Train/Validation lost and accuracy



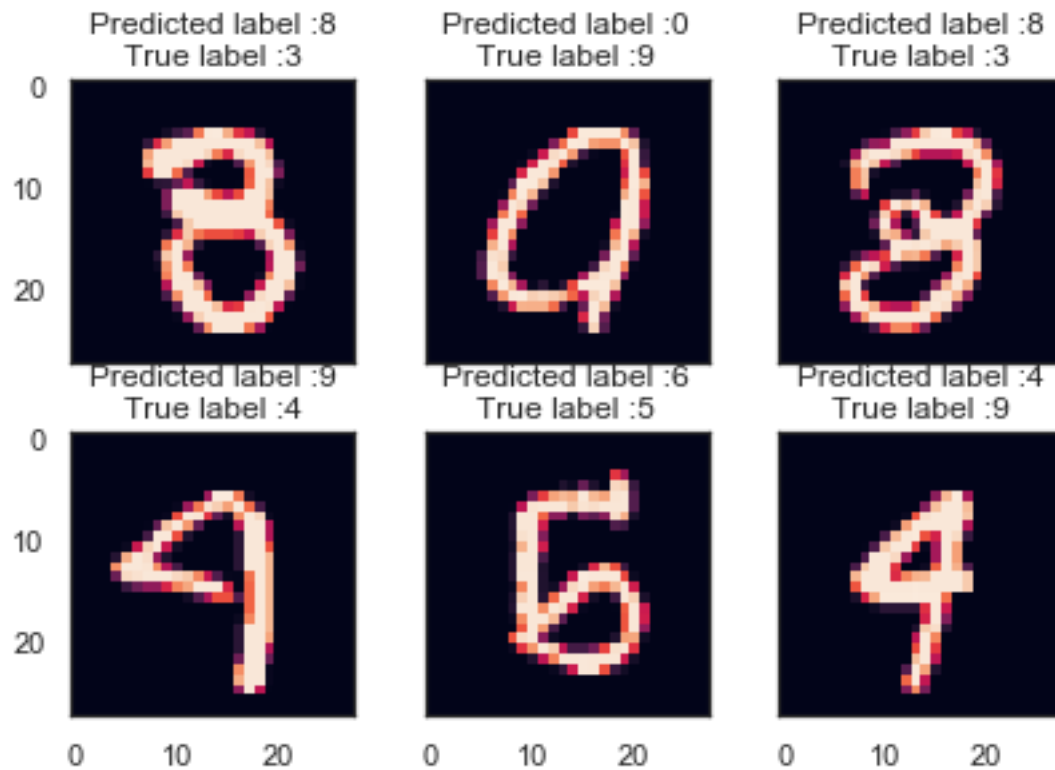
Result

- Confusion matrix



Result

- Top errors



Result

- Competition Ranking


— 636/2683 Top 24%

Entered

Sort by Grouped


All Categories


1 Active Competition



Digit Recognizer

Learn computer vision fundamentals with the famous MNIST data

Getting Started · Ongoing ·  tabular data, image data, multiclass classification, object identifica...



636/2683
Top 24%

No more competitions to show

Reference

- <https://www.kaggle.com/c/digit-recognizer>
- <https://www.kaggle.com/yassineghouzam/introduction-to-cnn-keras-0-997-top-6>
- <https://www.kaggle.com/leo223/keras-model-99-485-leaderboard-cnn/>
- https://en.wikipedia.org/wiki/MNIST_database
- <http://yann.lecun.com/exdb/mnist/>